

CS 188: Artificial Intelligence

Constraint Satisfaction Problems




Instructors: Josh Hug and Adam Janin

University of California, Berkeley

What is Search For?

- Assumptions about the world: a single agent, deterministic actions, fully observed state, discrete state space
- Planning: sequences of actions
 - The path to the goal is the important thing
 - Paths have various costs, depths
 - Heuristics give problem-specific guidance
- Identification: assignments to variables
 - The goal itself is important, not the path
 - All paths at the same depth (for some formulations)
 - CSPs are a special class of identification problems


Last week



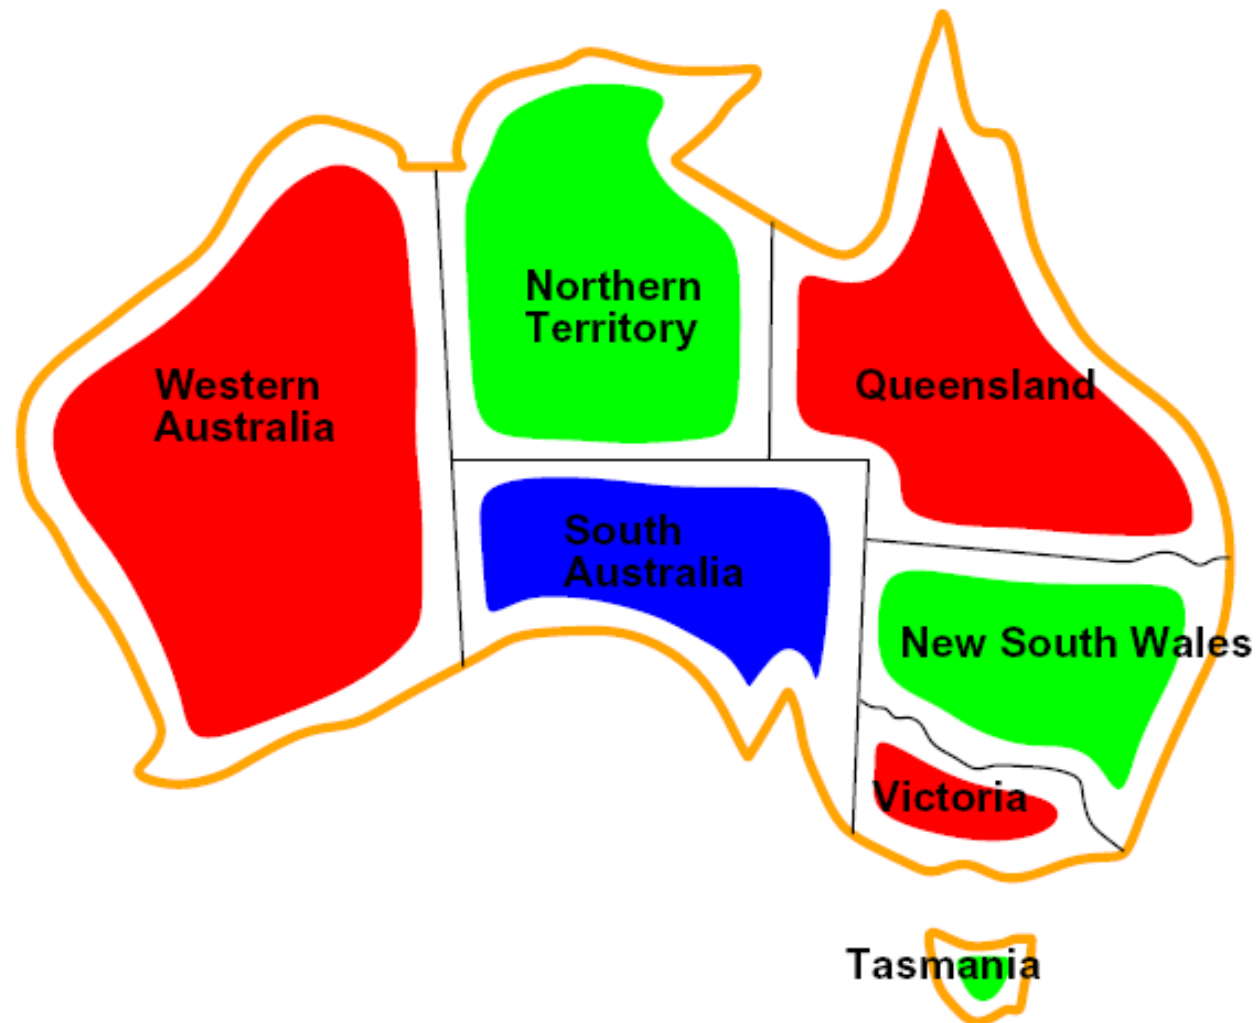
This week



Constraint Satisfaction Problems



CSP Examples



Example: Map Coloring



- Variables: WA, NT, Q, NSW, V, SA, T
- Domains: $D = \{\text{red, green, blue}\}$
- Constraints: adjacent regions must have different colors

Implicit: $WA \neq NT$

Explicit: $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$

- Solutions are assignments satisfying all constraints, e.g.:



$\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{green}\}$



Constraint Satisfaction Problems

- Standard search problems:
 - State is a “black box”: arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything
- Constraint satisfaction problems (CSPs):
 - A special subset of search problems
 - State is defined by **variables X_i** with values from a **domain D** (sometimes D depends on i)
 - Goal test is a **set of constraints** specifying allowable combinations of values for subsets of variables
- Simple example of a formal representation language
- Allows useful general-purpose CSP algorithms with more power than standard search algorithms

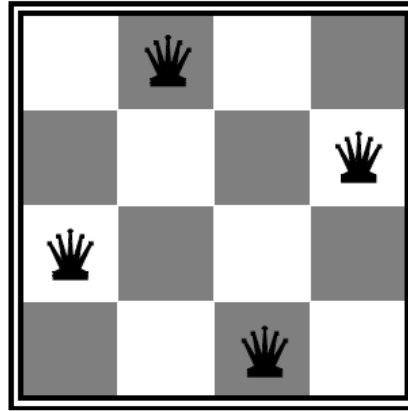
e.g., saw many possible formulations of Tower of Hanoi last week.




Example: N-Queens

■ Formulation 1:

- Variables: X_{ij}
- Domains: $\{0, 1\}$
- Constraints



16 variables for 4x4 case



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j+k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j-k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\sum_{i,j} X_{ij} = N$$

Example: N-Queens

- Formulation 2:

- Variables: Q_k


- Domains: $\{1, 2, 3, \dots, N\}$

- Constraints:

Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$


Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$ (6 possibilities)

...




4 variables for 4x4 case

Constraint Graphs



Constraint Graphs

- Binary CSP: each constraint relates (at most) two variables
- Binary constraint graph: nodes are variables, arcs show constraints
- General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!



Non-Binary CSPs, Example: Cryptarithmic

- Non-Binary CSP: constraints may relate more than two variables.

- Variables:

$F \ T \ U \ W \ R \ O \ X_1 \ X_2 \ X_3$

- Domains:

$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

- Constraints:

$\text{alldiff}(F, T, U, W, R, O)$

$O + O = R + 10 \cdot X_1$

...

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

$$\begin{array}{r} 7 \ 3 \ 4 \\ + \ 7 \ 3 \ 4 \\ \hline 1 \ 4 \ 6 \ 8 \end{array}$$


- In Cryptarithmic game:

- All letters represent a unique integer
- Goal is to find letters that are consistent with addition



Non-Binary CSPs

- Non-Binary CSP: may relate more than two variables.
- Non-binary constraint graph: nodes may be either variables, or special constraint nodes
 - Example: $\text{alldiff}(F, T, U, W, R, O)$ is represented by node at top



- Constraints:


$\text{alldiff}(F, T, U, W, R, O)$

$$O + O = R + 10 \cdot X_1$$

...

$$\begin{array}{r} T W O \\ + T W O \\ \hline F O U R \end{array}$$

Example: Sudoku



- Variables:
 - Each (open) square
- Domains:
 - $\{1,2,\dots,9\}$
- Constraints:

9-way alldiff for each column


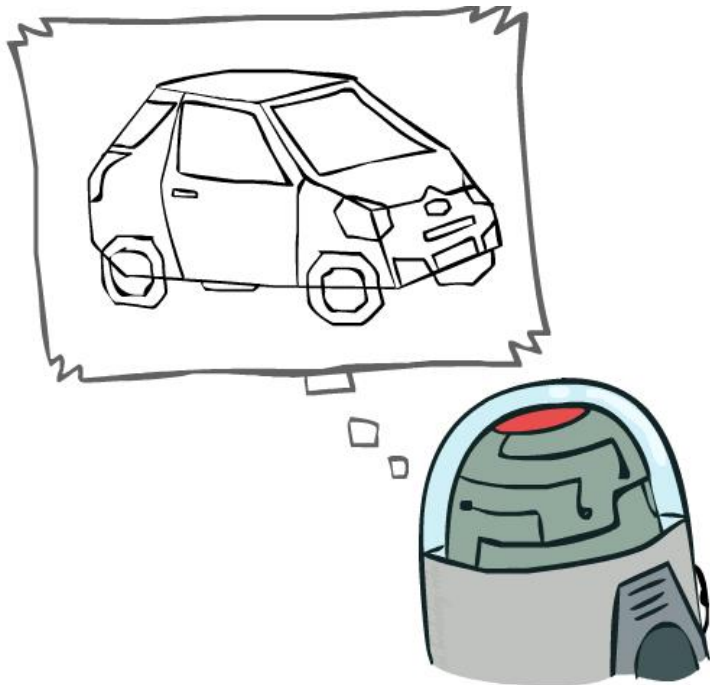
9-way alldiff for each row

9-way alldiff for each region

(other specifications are possible, e.g. can have a **bunch** of pairwise inequality constraints)

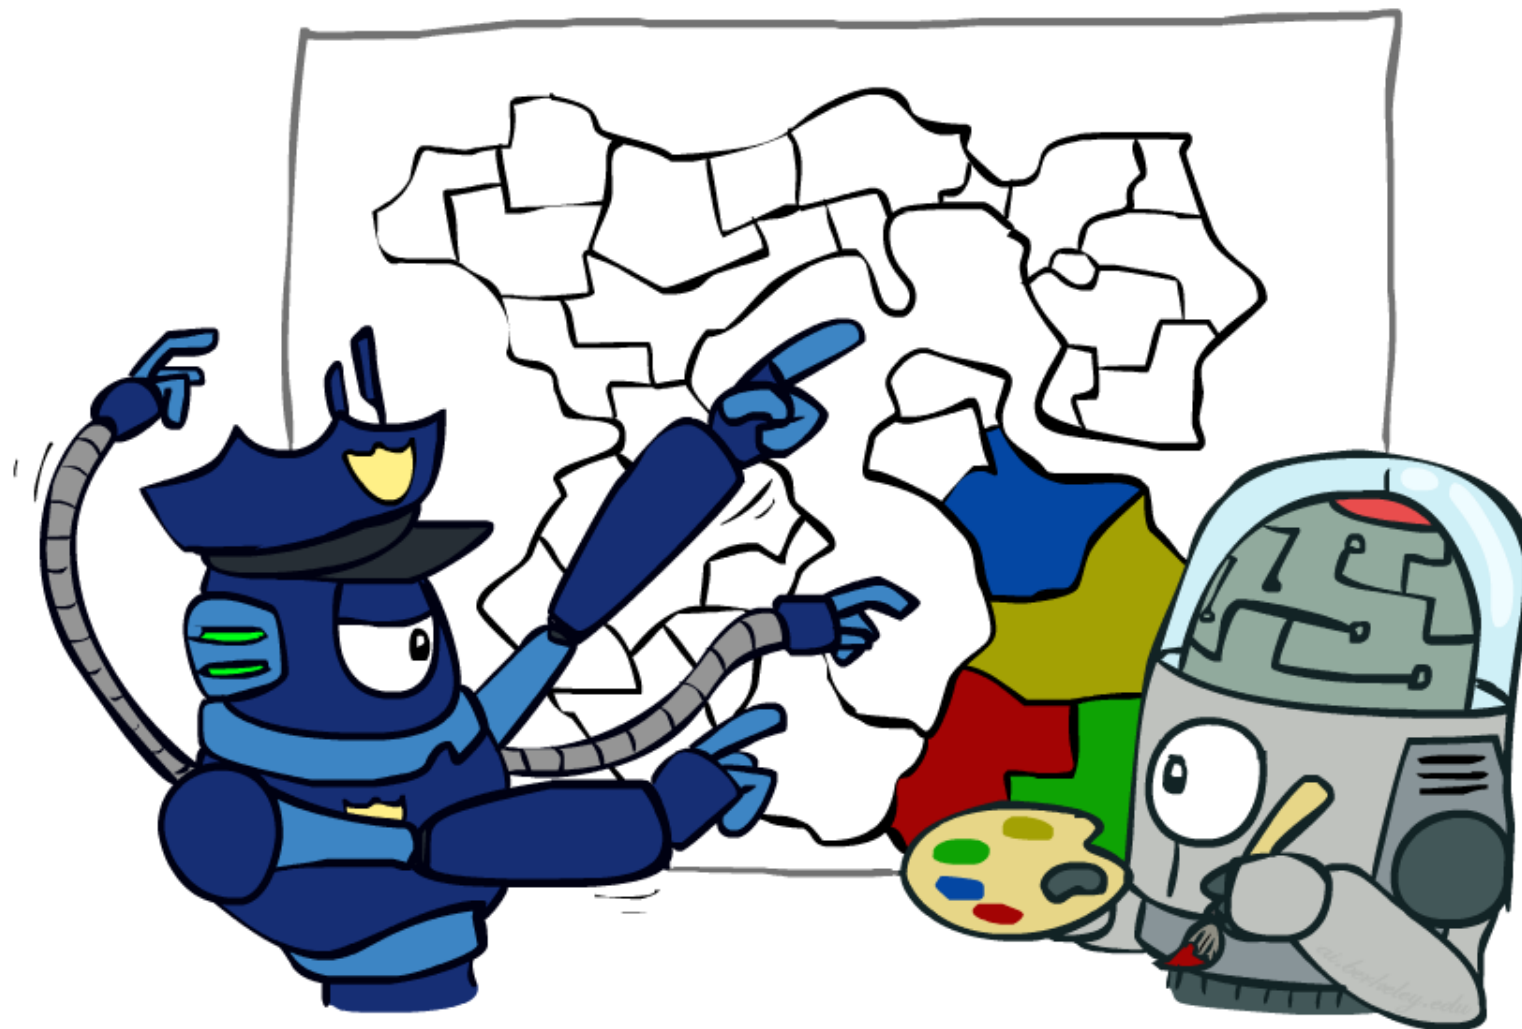
Example: The Waltz Algorithm

- The Waltz algorithm is for interpreting line drawings of solid polyhedra as 3D objects
- An early example of an AI computation posed as a CSP



- Approach:
 - Each intersection is a variable
 - Adjacent intersections impose constraints on each other
 - Solutions are physically realizable 3D interpretations

Varieties of CSPs and Constraints



Varieties of CSPs


- Discrete Variables

- Finite domains
 - Domain size d means $O(d^n)$ complete assignments
 - E.g., $d = 2$ includes 3-SAT Boolean satisfiability (NP-complete)
- Infinite domains (integers, strings, etc.)
 - E.g., job scheduling, variables are start/end times for each job
 - Linear constraints solvable, nonlinear undecidable

- Continuous variables

- E.g., start/end times for Hubble Telescope observations
- Linear constraints solvable in polynomial time by LP methods (see cs170 for a bit of this theory)

won't discuss in our course



Varieties of Constraints

- Varieties of Constraints

- Unary constraints involve a single variable (equivalent to reducing domains), e.g.:

$$SA \neq \text{green}$$

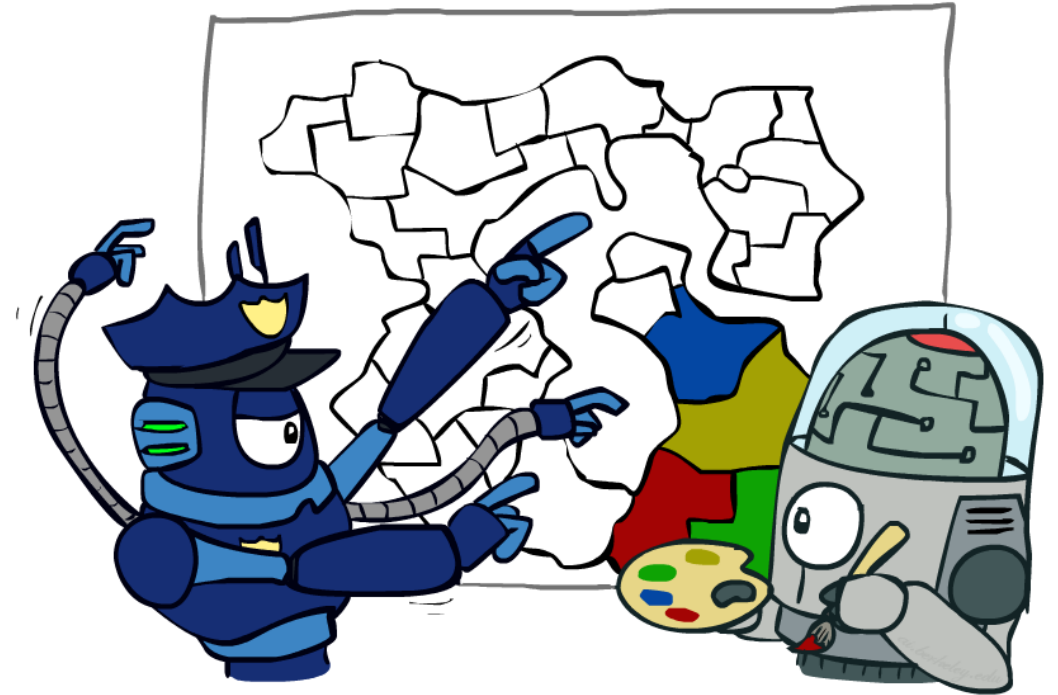
- Binary constraints involve pairs of variables, e.g.:

$$SA \neq WA$$

- Higher-order constraints involve 3 or more variables:
e.g., cryptarithmic column constraints

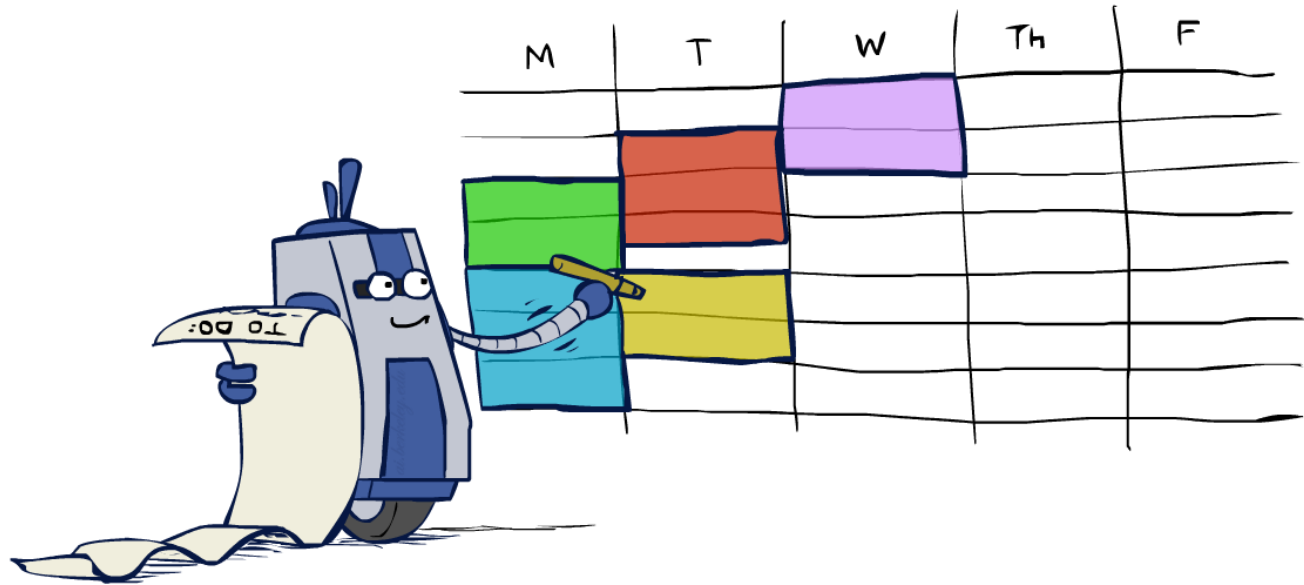
- Preferences (soft constraints):

- E.g., red is better than green
- Often representable by a cost for each variable assignment
- Gives constrained optimization problems
- (We'll ignore these until we get to Bayes' nets)



Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- Fault diagnosis
- ... lots more!




- Many real-world problems involve real-valued variables...

Solving CSPs



The Naive Formulation


- Standard search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
 - Initial state: the empty assignment, $\{\}$
 - Successor function: assign a value to an unassigned variable
 - Goal test: the current assignment is complete and satisfies all constraints
- We'll start with this straightforward, naïve approach, then improve it



Naive Search, i.e. Using DFS/BFS




- Challenge: Write out the first 3 states that BFS would expand.
- Challenge #2: Write out the first 3 states that DFS would expand.



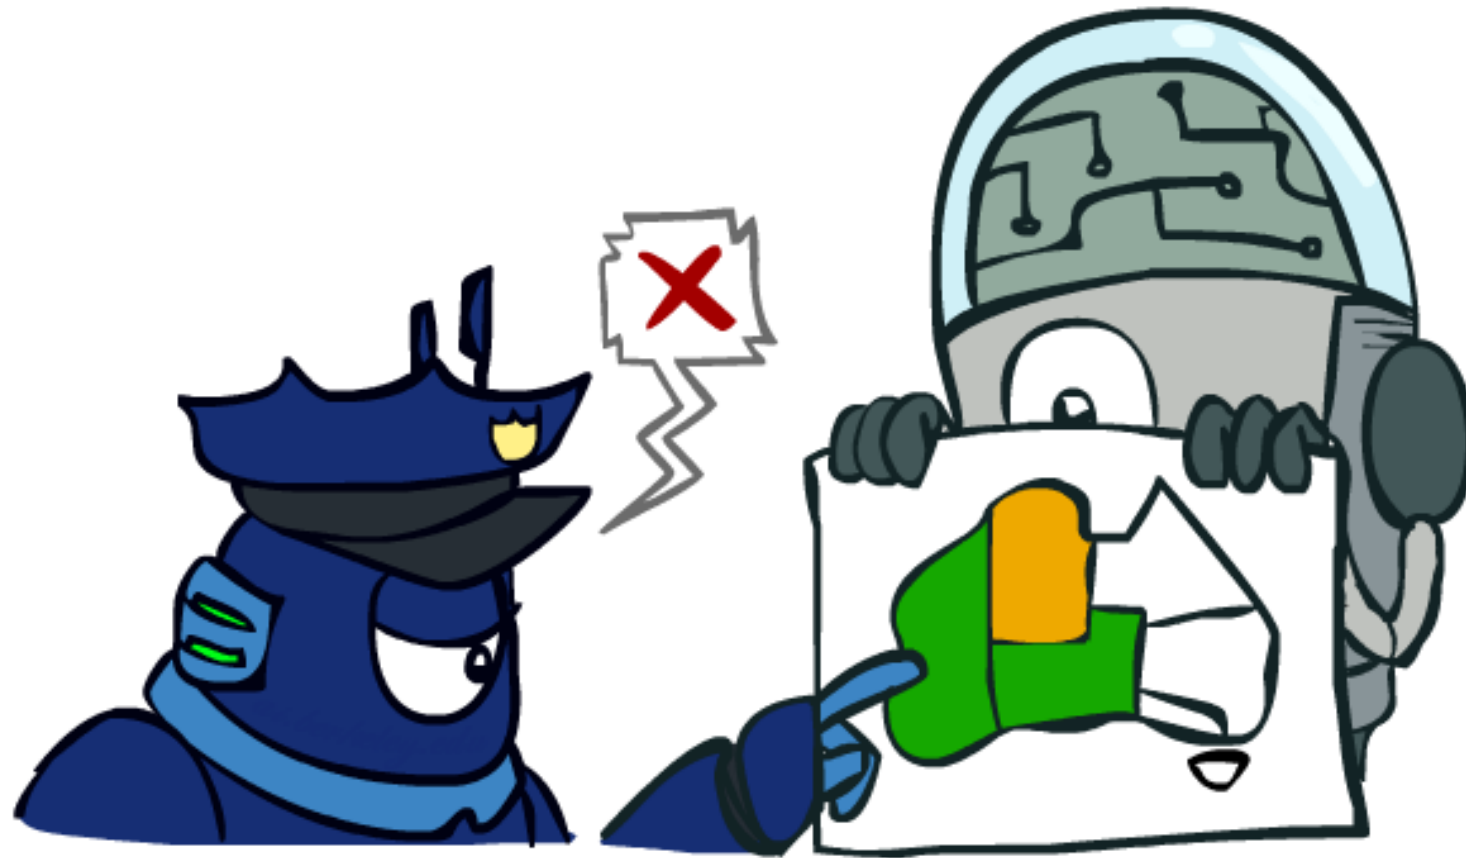
- Extra questions: Does the path matter? Where are the goal states? What are the problems with DFS & BFS?

Video of Demo Coloring -- DFS




[Demo: CSPs Demo]

Backtracking Search




Backtracking Search


- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
 - Variable assignments are commutative, so fix ordering
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
 - I.e. consider only values which do not conflict previous assignments
 - Might have to do some computation to check the constraint
 - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
 - I’ll use the term “CSP-backtracking” to avoid confusion.
- Can solve n-queens for $n \approx 25$



DFS (left) vs. CSP-Backtracking (Right)



Video of Demo Coloring – CSP-Backtracking



[Demo: CSPs Demo]

CSP-Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)


function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

- CSP-Backtracking = DFS + variable-ordering + fail-on-violation
- One optimization possibility: Pick “better” variable orderings and value orderings.

The Return of Question Goat



- Consider the partially completed CSP assignment. Assume the decisions about variable assignments are made bottom-up, left-to-right. Let X be the number of the decision that most obviously doomed the current assignment.
 - What is X ?
 - Bonus: How many decisions will be made before CSP-Backtracking search realizes its error?



Improving CSP-Backtracking


General-purpose ideas give huge gains in speed

- Ordering:

- Which variable should be assigned next?
- In what order should its values be tried?

- ? ■ Filtering: Can we detect inevitable failure early?

- Structure: Can we exploit the problem structure?




We'll start with filtering.

Filtering




Filtering: Forward Checking (Finding 8)

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: Cross off values that violate a constraint when added to the existing assignment




Coloring Demo – CSP-Backtracking with Forward Checking



Filtering: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:





WA	NT	Q	NSW	V	SA
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

- NT and SA cannot both be blue!
- Why didn't we detect this yet?
- Constraint propagation*: reason from constraint to constraint

Consistency of A Single Arc

- An arc $X \rightarrow Y$ is **consistent** iff for *every* x in the tail there is *some* y in the head which could be assigned without violating a constraint





Delete from the tail!

- Forward checking: Enforcing consistency of arcs pointing to each new assignment

Arc Consistency of an Entire CSP (Finding 4)

- A simple form of propagation makes sure **all** arcs are consistent:



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure earlier than forward checking (finding 4 vs. 8)
- Can be run as a preprocessor or after each assignment
- What's the downside of enforcing arc consistency?

*Remember: Delete
from the tail!*

AC-3: Enforcing Arc Consistency in a CSP

function **AC-3**(*csp*) **returns** the CSP, possibly with reduced domains

inputs: *csp*, a binary CSP with variables $\{X_1, X_2, \dots, X_n\}$

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$

if **REMOVE-INCONSISTENT-VALUES**(X_i, X_j) **then**

for each X_k **in** **NEIGHBORS**[X_i] **do**

add (X_k, X_i) to *queue*

tail **head**



function **REMOVE-INCONSISTENT-VALUES**(X_i, X_j) **returns** true iff succeeds

removed \leftarrow *false*

for each x **in** **DOMAIN**[X_i] **do**

if no value y in **DOMAIN**[X_j] allows (x, y) to satisfy the constraint $X_i \leftrightarrow X_j$


then delete x from **DOMAIN**[X_i]; *removed* \leftarrow *true*

return *removed*


The “delete
from trunk”
function.

- Runtime: $O(n^2d^3)$, can be reduced to $O(n^2d^2)$
- ... but detecting all possible future problems is NP-hard – why?



Video of Demo Arc Consistency – CSP Applet – n Queens



Limitations of Arc Consistency



- For the graphs to the right:
 - Are they arc consistent?
 - How many solutions are there to the CSP?



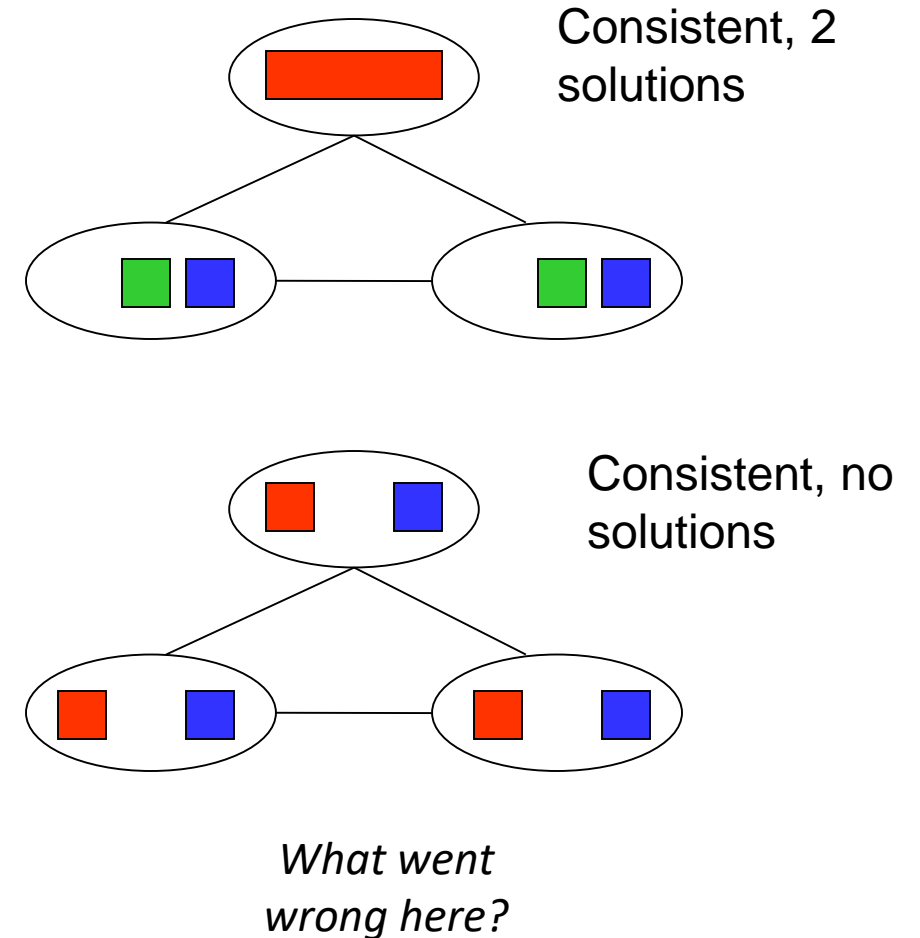
*What went
wrong here?*

[Demo: coloring -- forward checking]

[Demo: coloring -- arc consistency]

Limitations of Arc Consistency


- After enforcing arc consistency:
 - Can have one solution left
 - Can have multiple solutions left
 - Can have no solutions left (and not know it)
- Arc consistency still runs inside a CSP-backtracking search!




[Demo: coloring -- forward checking]

[Demo: coloring -- arc consistency]

Coloring Demo Video – CSP-Backtracking with Forward Checking on a Complex Graph



Coloring Demo Video – CSP-Backtracking with Arc Consistency – Complex Graph



Improving CSP-Backtracking

General-purpose ideas give huge gains in speed


? ■ Ordering:

- Which variable should be assigned next?
- In what order should its values be tried?

- ✓ Filtering: Can we detect inevitable failure early?

- Yes! Use forward checking or arc consistency (e.g. AC-3).
- AC-3 more expensive, but catches problems earlier.

- Structure: Can we exploit the problem structure?



Ordering




Ordering: Minimum Remaining Values

- Variable Ordering: Minimum remaining values (MRV):
 - Choose the variable with the fewest legal values remaining in its domain




- Why min rather than max?
- Also called “most constrained variable”
- “Fail-fast” ordering



Ordering: Least Constraining Value


- Value Ordering: Least Constraining Value

- After picking a variable, which value should we pick? Choose the *least constraining value*
- i.e., the one that rules out the fewest values in the remaining variables
- Note that it may take some computation to determine this! (E.g., rerunning filtering)



- Why least rather than most?

- Don't have to try every value! Use the one most likely to work.



- Combining these ordering ideas makes 1000 queens feasible

Demo: Coloring – CSP-Backtracking + Forward Checking + Ordering
