

CS 188: Artificial Intelligence

Constraint Satisfaction Problems II

Instructors: Adam Janin and Josh Hug

University of California, Berkeley



Today

- Efficient Solution of CSPs
 - A few more improvements to CSP-Backtracking.
 - An alternate approach: Iterative Improvement.


- Local Search



Reminder: CSPs


■ CSPs:

- Variables
- Domains
- Constraints
 - Implicit (provide code to compute)
 - Explicit (provide a list of the legal tuples)
 - Unary / Binary / N-ary



■ Goals:

- Here: find any solution
- Also: find all, find best, etc.



Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Improving CSP-Backtracking


General-purpose ideas give huge gains in speed

- ✓ Ordering: Can we pick better orderings?
 - Yes! Can use most constrained **variable** (MCV) and least constraining **value** (LCV).
- ? ■ Filtering: Can we detect inevitable failure early?
 - ✓ Yes! Use forward checking or arc consistency (e.g. AC-3).
 - AC-3 more expensive, but catches problems earlier.
 - ? ■ Will discuss new idea of “K-consistency” briefly today.
- ? ■ Structure: Can we exploit the problem structure?

Note: Overall problem is still NP-hard.




Arc Consistency and Beyond



An arc $X \rightarrow Y$ is **consistent** iff **for every** x in the tail there is **some** y in the head which could be assigned without violating a constraint


Arc Consistency Practice/Review

- Variables: X, Y, Z
- Domains: {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16}
- Constraints:
 - X is either the square or the square root of Y.
 - Y is either the square or the square root of Z.
 - X is less than or equal to Z.




- Some sample solutions:
 - X, Y, Z = {1, 1, 1}
 - X, Y, Z = {2, 4, 16}
 - X, Y, Z = {9, 3, 9}
- Problem 1: If we run AC-3 on this constraint graph, what are the possible values for X, Y, Z?
- Problem 2: If we run AC-3, then assign X = 16, then run arc-consistency again, what are the possible values for Y and Z?


Arc Consistency Practice




- Variables: X, Y, Z
- Domains: {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16}
- Problem 1: If we run AC-3 on this constraint graph, what are the possible values for X, Y, Z?



Arc Consistency Practice





- Variables: X, Y, Z
- Domains: {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16}
- Problem 2: If we run AC-3, assign X = 16 and run AC-3 again, what are the possible values for Y and Z?



Reminder: If you delete from a node, must re-enqueue all arcs pointing at that node!


Limitations of Arc Consistency

- After enforcing arc consistency:
 - Can have one solution left
 - Can have multiple solutions left
 - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!






*What went
wrong here?*

K-Consistency



K-Consistency


- Increasing degrees of consistency
 - 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
 - 2-Consistency (Arc Consistency)**: For each pair of nodes, any consistent assignment to one can be extended to the other
 - K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the kth node
- Higher k more expensive to compute
- (You need to know the k=2 case: **arc consistency**)



This graph is 2-consistent, but not 3-consistent.

Strong K-Consistency

- Strong k-consistency: also k-1, k-2, ... 1 consistent
- Claim: strong n-consistency means we can solve without backtracking!
- Why?
 - Choose any assignment to any variable
 - Choose a new variable
 - By 2-consistency, there is a choice consistent with the first
 - Choose a new variable
 - By 3-consistency, there is a choice consistent with the first 2
 - ...
- Lots of middle ground between arc consistency (a.k.a. 2-consistency) and n-consistency!
 - Example, k = 3 (path consistency) helps avoid situations like that to the right.
 - Won't discuss algorithms for $k > 2$. If you're curious see [this link](#).



Improving CSP-Backtracking


General-purpose ideas give huge gains in speed

- ✓ Ordering: Can we pick better orderings?
 - Yes! Can use most constrained **variable** (MCV) and least constraining **value** (LCV).
- ✓ Filtering: Can we detect inevitable failure early?
 - Yes! Use forward checking, **arc consistency**, or K-consistency.
 - K-consistency > **arc consistency** > forward checking in both cost and quality of filtering (assuming $k > 2$).
- ? ■ Structure: Can we exploit the problem structure?

Note: Overall problem is still NP-hard.




Structure




Problem Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into n/c subproblems of only c variables:
 - Worst-case solution cost is $O\left(\frac{n}{c} d^c\right)$, linear in n
 - E.g., $n = 80$ variables, $d = 2$, $c = 20$, broken into 4 subproblems
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



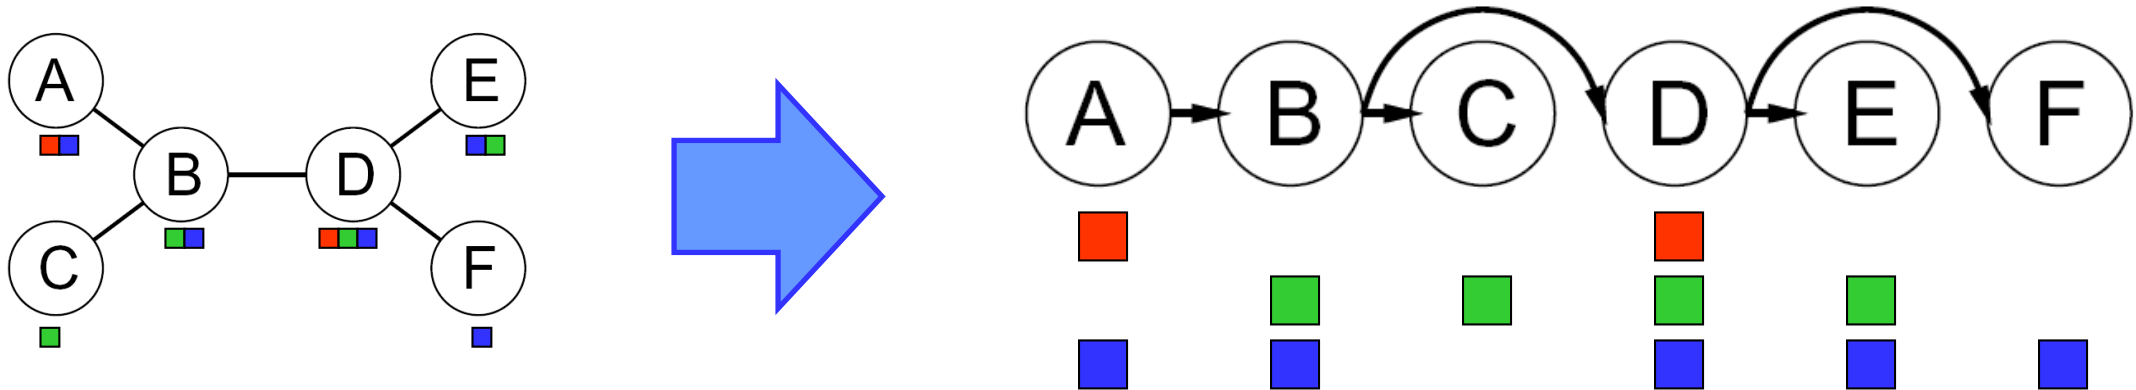
Tree-Structured CSPs




- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(n d^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$
- This property also applies to probabilistic reasoning (later): an example of the relation between syntactic restrictions and the complexity of reasoning

Tree-Structured CSPs

- Algorithm for tree-structured CSPs:
 - Order: Choose a root variable, order variables so that parents precede children

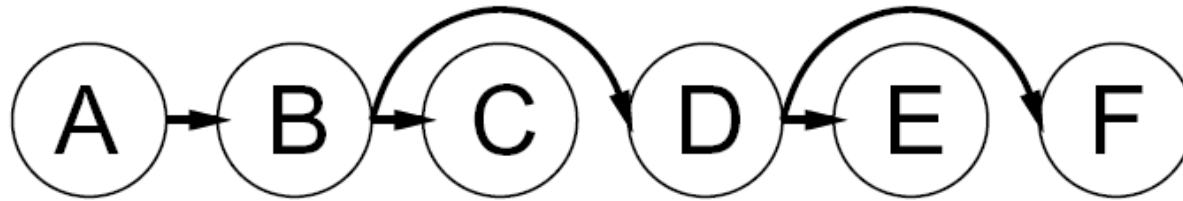


- Remove backward: For $i = n : 2$, apply $\text{RemoveInconsistent}(\text{Parent}(X_i), X_i)$
 - Assign forward: For $i = 1 : n$, assign X_i consistently with $\text{Parent}(X_i)$
- Runtime: $O(n d^2)$ (why?)




Tree-Structured CSPs (Tricky Slide. Revisit Later!)

- Claim 1: After backward pass, all parent-to-child arcs are 2-consistent
- Proof: Each $X \rightarrow Y$ was made 2-consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)




- Claim 2: If parent-to-child arcs are consistent, forward assignment will not backtrack
- Proof: Induction on position, e.g. $A \rightarrow B$ 2-consistent means **any** choice for A is safe for B .
- Why doesn't this algorithm work with cycles in the constraint graph?
 - Arc-consistency doesn't guarantee safety for multiple parents! Induction on position fails.
- Note: we'll see this basic idea again with Bayes' nets

Improving Structure



Nearly Tree-Structured CSPs



- Conditioning: instantiate a variable, prune its neighbors' domains
- Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size c gives runtime $O(d^c (n - c) d^2)$, very fast for small c


Cutset Conditioning

Choose a cutset


Instantiate the cutset
(all possible ways)

Compute residual CSP
for each assignment


Solve the residual CSPs
(tree structured)



Finding the Minimum Cutset



- Find the smallest cutset for the graph below.
 - Set of nodes which, when removed, results in a tree.




Improving CSP-Backtracking

General-purpose ideas give huge gains in speed

- ✓ **Ordering:** Can we pick better orderings?
 - Yes! Can use most constrained **variable** (MCV) and least constraining **value** (LCV)
- ✓ **Filtering:** Can we detect inevitable failure early?
 - Yes! Use forward checking, arc consistency, or K-consistency
 - K-consistency > arc consistency > forward checking in both cost and quality of filtering
- ✓ **Structure:** Can we exploit the problem structure?
 - Tree CSPs: Use tree-algorithm
 - Almost tree CSPs: Use cutset conditioning
 - Wanna get fancy? Try Tree Decomposition




Alternate Approach: Iterative Improvement




CSP-Backtracking works, but there are other approaches. Iterative Improvement is one such approach.

Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - Take an assignment with unsatisfied constraints
 - Operators *reassign* variable values
 - No fringe! Live on the edge.
- Algorithm: While not solved,
 - Variable selection: randomly select any conflicted variable
 - Value selection: min-conflicts heuristic:
 - Choose a value that violates the fewest constraints
 - i.e., hill climb with $h(n)$ = total number of violated constraints
 - i.e., don't change values if all the other possibilities are actually worse




Example: 4-Queens




- States: 4 queens in 4 columns ($4^4 = 256$ states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation: $c(n)$ = number of attacks

Video of Demo Iterative Improvement – n Queens



Video of Demo Iterative Improvement – Coloring



Performance of Min-Conflicts

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$







Figure via Stuart Russell

Min-Conflicts and 5-Queens




- For the 5-Queens board shown:
 - How many conflicts are there?
 - What is the best move that Min-Conflicts can make (in terms of # of conflicts)?
 - Recall: Can only move a single queen within a single row.
 - How many conflicts remain after taking the best move?




Min-Conflicts, 5-Queens, Sideways Moves

- For the 5-Queens board shown:
 - How many conflicts are there?
 - 2
 - What is the best move that Min-Conflicts can make (in terms of # of conflicts)?
 - There are many possibilities, e.g. moving Queen 2 to Column 2.
 - How many conflicts remain after taking the best move?
 - 2
- A move that does not improve things is known as a “sideways” move.
 - This is in contrast to an “uphill” move.



Summary: CSPs

- CSPs are a special kind of search problem:
 - States are partial assignments
 - Goal test defined by constraints
- Basic solution: CSP-backtracking search
- Speed-ups:
 - Ordering
 - Filtering
 - Structure
- Alternate solution: Iterative Min-conflicts. Often effective in practice, particularly for huge problems.



Local Search




Local Search

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
 - Examples: DFS, BFS, UCS, A*, CSP-Backtracking.

- Local search: improve a single option until you can't make it better (no fringe!)
 - Example: Min-Conflicts

- New successor function: local changes




- Generally much faster and more memory efficient (but incomplete and suboptimal)

Hill Climbing



- Simple, general idea:
 - Start wherever
 - Repeat: move to the best neighboring state
 - If **no neighbors better** than current, quit
- What's bad about this approach?
 - Complete?
 - Optimal?
- What's good about it?



Hill Climbing Diagram



Hill-Climbing Quiz



Starting from X, where do you end up ?

Starting from Y, where do you end up ?


Starting from Z, where do you end up ?

Simulated Annealing

- Idea: Escape local maxima by allowing downhill moves ([demo](#))
 - But make them rarer as time goes on


```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                   next, a node
                   T, a “temperature” controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E / T}$ 
```




Simulated Annealing

- Theoretical guarantee:
 - If T decreased slowly enough, will converge to optimal state!
- Is this an interesting guarantee?
 - Kinda. Have to run algorithm forever to get guaranteed optimum.
- Sounds like magic, but reality is reality:
 - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
 - People think hard about *ridge operators* which let you jump around the space in better ways



Random Restarts


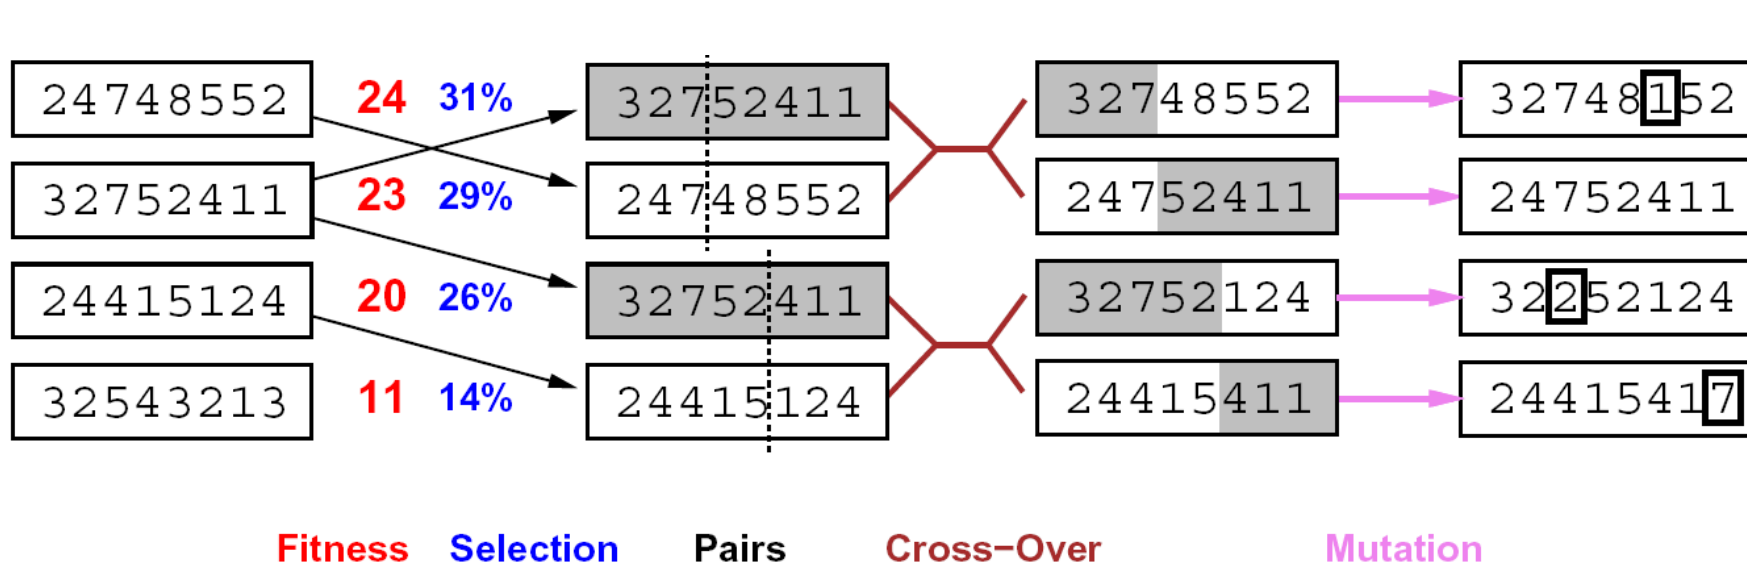
- Vanilla hill climbing and simulated annealing can be parallelized:
 - After running for a while, restart from a new random starting point.
 - Restart conditions include:
 - After running for Z total steps.
 - After spending Z units of time on a shoulder.
 - Anything else you can dream.
 - Can run many random searches in parallel.



Number of constraints obeyed


4th search performs random restart.

Genetic Algorithms



- Genetic algorithms use a natural selection metaphor
 - Keep best N hypotheses at each step (selection) based on a fitness function
 - Also have pairwise crossover operators, with optional mutation to give variety
 - Instead of taking a step in your landscape, take a huge random leap
- Possibly the most misunderstood, misapplied (and even maligned) technique around

Example: N-Queens



- Why does crossover make sense here?
- When wouldn't it make sense?
- What would mutation be?
- What would a good fitness function be?

Next Time: Adversarial Search!

- Example: Picking a move in chess by considering potential opponent's moves.

