


Announcements

- Homework 4: MDPs (today's topic)
 - Due Monday 9/26 at 11:59pm.
- Project 2: Multi-Agent Pacman
 - Has been released, due Friday 9/30 at 5:00pm.
- Contest 2, due at a time TBA, but soon after project 2.
 - Small tweak from contest 1.
- Survey on how we're doing so far has been released.
 - Due Saturday 9/24 at 11:59 PM.
 - +1 project point added to total (equivalent to one contest).
- Midterm: Oct 6
 - Time to start studying.
 - Recommended approach:
 - Work through problems alone.
 - Get together in group of 3-6, and whiteboard attempted solutions.
 - Interrupt each other!
 - Fill out midterm conflict form ASAP. Makeup at 8 AM.

CS 188: Artificial Intelligence


Markov Decision Processes II



Instructors: Dan Klein and Pieter Abbeel --- University of California, Berkeley

Example: Grid World

- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- Noisy movement: actions do not always go as planned
 - 80% chance: agent goes the way it wants to go (e.g. the action North takes the agent North)
 - 10% chance: agent steps left (e.g. North, but goes West)
 - 10% chance: agent steps right (e.g. North, but goes East)
 - If there is a wall in the direction the agent would have been taken, the agent stays put
 - 0% chance to go backwards
- The agent receives rewards each time step
 - Small "living" reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards




Recap: MDPs

- Markov decision processes:

- States S
- Actions A
- Transitions $P(s' | s, a)$ (or $T(s, a, s')$)
- Rewards $R(s, a, s')$ (and discount γ)
- Start state s_0


- Quantities:

- Policy = map of states to actions
- Utility = sum of discounted rewards
- Values = expected future utility from a state (max node)
- Q-Values = expected future utility from a q-state (chance node)



Optimal Quantities

- The value (utility) of a state s :
 $V^*(s)$ = expected utility starting in s and acting optimally
- The value (utility) of a q -state (s,a) :
 $Q^*(s,a)$ = expected utility starting out having taken action a from state s and (thereafter) acting optimally
- The optimal policy:
 $\pi^*(s)$ = optimal action from state s




Gridworld Display




VALUES AFTER 100 ITERATIONS

Gridworld: Q^*



The Bellman Equations



How to be optimal:

Step 1: Take correct first action

Step 2: Keep being optimal

The Bellman Equations


- Definition of “optimal utility” via expectimax recurrence gives a simple one-step lookahead relationship amongst optimal utility values

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- These are the Bellman equations, and they characterize optimal values in a way we'll use over and over



Value Iteration


- Bellman equations **characterize** the optimal values:

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Value iteration **computes** them:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Value iteration is just a fixed point solution method
 - ... though the V_k vectors are also interpretable as time-limited values



Value Iteration (Alternate View)

- **Bellman Equation** gives us a way to **characterize** whether a vector of utilities is “correct”, i.e the vector is equal to the expected utilities if we play optimally.

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$V(s)$

2.375 1.375 0

If we plug in our $V(\text{cool})$, $V(\text{warm})$, and $V(\text{OH})$, and the equation is satisfied, then our vector V is the expected utility if we play optimally.

- Can use **value iteration** to iteratively transform an initially all zero vector into the “correct” vector. Or in other words, lets us **compute** the “correct” vector.

0 0 0

$V_0(s)$

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$


2 1 0

$V_1(s)$




...

Convergence


- How do we know the V_k vectors are going to converge?
- Case 1: If the tree has maximum depth M , then V_M holds the actual untruncated values
- Case 2: If the discount is less than 1
 - Sketch: For any state V_k and V_{k+1} can be viewed as depth $k+1$ expectimax results in nearly identical search trees
 - The difference is that on the bottom layer, V_{k+1} has actual rewards while V_k has zeros
 - That last layer is at best all R_{MAX}
 - It is at worst R_{MIN}
 - But everything is discounted by γ^k that far out
 - So V_k and V_{k+1} are at most $\gamma^k \max |R|$ different
 - So as k increases, the values converge



Issues with Value Iteration

			
V_2	2.3 (F)	1.3 (S)	0
V_1	2 (F)	1 (S)	0
V_0	0	0	0


Assuming discount, $\gamma = 0.2$.




$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Issue #1: Slow, takes $O(S^2A)$ time.
- Issue #2: Clearly terrible choices keep being tested.
 - e.g. no need to **ever** test the choice: (Warm, Fast).
- Issue #3: Policy converges long before the values.
 - If using value iteration to find best policy, don't actually care about values.

Policy Methods




Policy Evaluation




Fixed Policies

Do the optimal action



Do what π says to do




- Expectimax trees max over all actions to compute the optimal values
- If we fixed some policy $\pi(s)$, then the tree would be simpler – only one action per state
 - ... though the tree's value would depend on which policy we fixed

Invent the Bellman Equation (Round 2)



- Another basic operation: compute the utility of a state s under a fixed (generally non-optimal) policy
- Define the utility of a state s , under a fixed policy π :
 $V^\pi(s)$ = expected total discounted rewards starting in s and following π
- Challenge: Write the Bellman Equation for $V^\pi(s)$



Hint:
$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


Playing optimally

$$V^\pi(s) =$$

Playing with π

Invent the Bellman Equation (Round 2)

- Another basic operation: compute the utility of a state s under a fixed (generally non-optimal) policy
- Define the utility of a state s , under a fixed policy π :
 $V^\pi(s)$ = expected total discounted rewards starting in s and following π
- Challenge: Write the Bellman Equation for $V^\pi(s)$



Hint:
$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


Playing optimally

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$


Playing with π

Example: Policy Evaluation

Always Go Right



Always Go Forward



Example: Policy Evaluation

Always Go Right



Always Go Forward




Policy Evaluation

- How do we calculate the V 's for a fixed policy π ?
- Idea 1: Turn recursive Bellman equations into updates (like value iteration)

$$V_0^\pi(s) = 0$$

$$V_{k+1}^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$

- Efficiency: $O(S^2)$ per iteration
- Idea 2: Without the maxes, the Bellman equations are just a linear system
 - Solve with Matlab (or your favorite linear system solver)



Policy Evaluation Exercise



Evaluate the policy $\pi(s) = \text{Slow}$




V_2

V_1

V_0

000

Assume discount $\gamma = 0.2$.



$$V_{k+1}^{\pi}(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$

Bonus question, what is $V^{\pi}(s)$?

Policy Evaluation Exercise



Evaluate the policy $\pi(s) = \text{Slow}$



V_2

1.2 (S)	1.2 (S)	0
---------	---------	---


V_1

1 (S)	1 (S)	0
-------	-------	---

V_0

0	0	0
---	---	---

Assume discount $\gamma = 0.2$.




$$V_{k+1}^{\pi}(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^{\pi}(s')]$$


Bonus: Under the given policy, $V^{\pi}(\text{cool})$ and $V^{\pi}(\text{warm})$ converge to 1.25.

Value Iteration vs. Policy Evaluation

- Value iteration lets us **compute** the expected utility if we play optimally.




- Policy evaluation lets us **compute** the expected utility if we play using some policy $\pi(s)$.




Value Iteration vs. Policy Evaluation



- Value iteration lets us **compute** the expected utility if we play optimally.



- Policy evaluation lets us **compute** the expected utility if we play using some policy $\pi(s)$.




- Let **X(s)** be the vector resulting from running policy evaluation on the **optimal policy**.
- Let **Y(s)** be the vector resulting from running policy evaluation on a **policy** which always selects the “leftmost” action, i.e. first action on the list of possible actions.
- Let **Z(s)** be the vector resulting from running **value iteration**.


Assuming we run them all to convergence, how are the values of X, Y, and Z related?

Value Iteration vs. Policy Evaluation

- Value iteration lets us **compute** the expected utility if we play optimally.




- Policy evaluation lets us **compute** the expected utility if we play using some policy $\pi(s)$.



- Let **X(s)** be the vector resulting from running **policy evaluation** on the **optimal policy**.
- Let **Y(s)** be the vector resulting from running **policy evaluation** on a **policy** which always selects the “leftmost” action, i.e. first action on the list of possible actions.
- Let **Z(s)** be the vector resulting from running **value iteration**.

Answer: X and Z are equal. All values in Y are less than or equal to X.


Policy Extraction



Computing Actions from Q-Values

- Let's imagine we have the optimal q-values:
- What is the optimal policy?
 - Completely trivial to decide!


$$\pi^*(s) = \arg \max_a Q^*(s, a)$$



- Important lesson: actions are easier to select from q-values than values!

Computing Actions from Values

- Let's imagine we have the optimal values $V^*(s)$
- What is the optimal policy, i.e. how should we act?
 - It's not obvious!



???



Computing Actions from Values

- Let's imagine we have the optimal values $V^*(s)$
- What is the optimal policy?
 - It's not obvious!
- We need to do a mini-expectimax (one step)
 - In other words, compute q-values, pick action that goes with maximum for each state.


0.95	0.96	0.98	1.00
0.94		0.89	-1.00
0.92	0.91	0.90	0.80

$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


- This is called **policy extraction**, since it gets the policy implied by the values
 - Important observation: actions are easier to select from q-values than values!

Value Iteration, Policy Evaluation, Policy Extraction


- Value iteration lets us **compute** the expected utility if we play optimally.




- Policy evaluation lets us **compute** the expected utility if we play using some policy $\pi(s)$.



- Policy extraction lets us **compute** the optimal policy from an optimal vector $V^*(s)$.




Policy Iteration




Problems with Value Iteration


- Value iteration repeats the Bellman updates:



- Problem 1: It's slow – $O(S^2A)$ per iteration
- Problem 2: The best action at each state rarely changes
- Problem 3: The policy often converges long before the values



$k=0$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=1



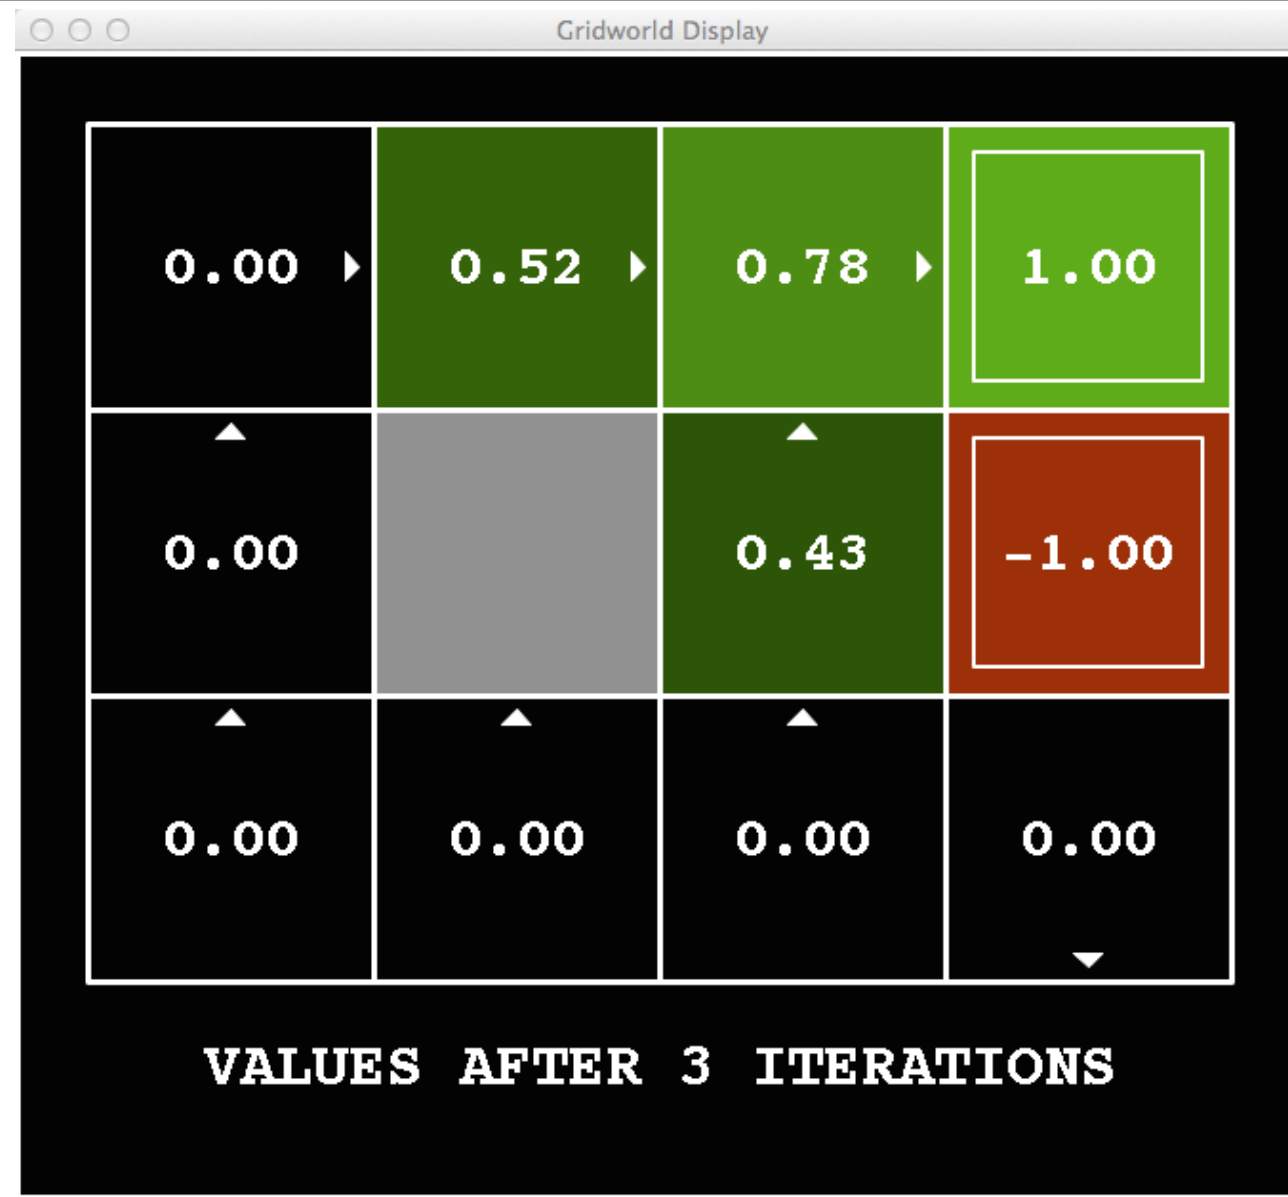
Noise = 0.2
Discount = 0.9
Living reward = 0

k=2



Noise = 0.2
Discount = 0.9
Living reward = 0

k=3



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=4$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=5



Noise = 0.2
Discount = 0.9
Living reward = 0

k=6



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=7$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=8



Noise = 0.2
Discount = 0.9
Living reward = 0

k=9



Noise = 0.2
Discount = 0.9
Living reward = 0

k=10



Noise = 0.2
Discount = 0.9
Living reward = 0

k=11



Noise = 0.2
Discount = 0.9
Living reward = 0

k=12



Noise = 0.2
Discount = 0.9
Living reward = 0

k=100



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

- Alternative approach for optimal values:
 - **Step 1: Policy evaluation:** calculate utilities for some fixed policy (not optimal utilities!) until convergence
 - **Step 2: Policy improvement:** update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
 - Repeat steps until policy converges
- This is **policy iteration**
 - It's still optimal!
 - Can converge (much) faster under some conditions

Policy Iteration

- Evaluation: For fixed current policy π_i , find values with **policy evaluation**:
 - Iterate until values converge:


$$V_{k+1}^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

- Improvement: For fixed values, get a better policy using **policy extraction**
 - Do a one-step look ahead using expectimax.
 - In other words, compute q-values, pick action that goes with maximum for each state.


$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

Policy Evaluation, Extraction, and Iteration

- Policy **evaluation** lets us **compute** the expected utility if we play using some policy $\pi(s)$.




- Policy **extraction** lets us **compute** the optimal policy from an optimal vector $V^*(s)$.




- Policy iteration is a call to **evaluate**, followed by a call to **extract**.

- Given a policy π_i , policy iteration outputs a better policy π_{i+1} .



More Thorough Example

- Come up with an **arbitrary initial policy** π_0 , say $\pi_0(s) = [\text{slow}, \text{slow}, \emptyset]$.
- Compute the expected utilities for this policy if we use policy π_0 using **policy evaluation**.
 - After converging (possibly many many iterations), this yields a vector $V^{\pi_0}(s)$.




More Thorough Example

- Come up with an **arbitrary initial policy** π_0 , say $\pi_0(s) = [\text{slow}, \text{slow}, \emptyset]$.
- Compute the expected utilities for this policy if we use policy π_0 using **policy evaluation**.
 - After converging (possibly many many iterations), this yields a vector $V^{\pi_0}(s)$.
 - The expected utility of the policy $[\text{slow}, \text{slow}, \emptyset]$ is $[1.25, 1.25, 0]$
- Compute a better policy π_1 using **policy iteration**.
 - In other words, compute q-values for each possible choice, pick new best.

1.25 1.25 0

$V_{182}^{\pi_0}(s)$



- Can repeat this process to get even better policies π_2, π_3 , etc.



Value Iteration vs. Policy Iteration

- Both value iteration and policy iteration compute the same thing (all optimal values)
- In value iteration (previous lecture):
 - Every iteration updates both the values and (implicitly) the policy
 - We don't track the policy, but taking the max over actions implicitly recomputes it
- In policy iteration (this lecture):
 - We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
 - The new policy will be better (or we're done)
- Both are dynamic programs for solving MDPs

Summary: MDP Algorithms


- So you want to....
 - Compute optimal values: use **value iteration** or **policy iteration**
 - Compute values for a particular policy: use policy evaluation
 - Turn your values into a policy: use policy extraction (one-step lookahead)
- These all look the same!
 - They basically are – they are all variations of Bellman updates
 - They all use one-step lookahead expectimax fragments
 - They differ only in whether we plug in a fixed policy or max over actions

Double Bandits



Double-Bandit MDP

- Actions: *Blue*, *Red*
- States: *Win*, Lose




No discount
100 time steps
Both states have the same value

Offline Planning


- Solving MDPs is offline planning
 - You determine all quantities through computation
 - You need to know the details of the MDP
 - You do not actually play the game!

No discount
100 time steps
Both states have the same value

	Value
Play Red	150
Play Blue	100




Let's Play!




\$2 \$2 \$0 \$2 \$2
\$2 \$2 \$0 \$0 \$0

Online Planning

- Rules changed! Red's win chance is different.



Let's Play!



\$0 \$0 \$0 \$2 \$0
\$2 \$0 \$0 \$0 \$0

What Just Happened?

- That wasn't planning, it was learning!
 - Specifically, reinforcement learning
 - There was an MDP, but you couldn't solve it with just computation
 - You needed to actually act to figure it out
- Important ideas in reinforcement learning that came up
 - Exploration: you have to try unknown actions to get information
 - Exploitation: eventually, you have to use what you know
 - Regret: even if you learn intelligently, you make mistakes
 - Sampling: because of chance, you have to try things repeatedly
 - Difficulty: learning can be much harder than solving a known MDP



Next Time: Reinforcement Learning!
