


# Announcements

---

- No Lecture this Thursday
  - Study instead
- Midterm this Thursday 10/4 at 7 PM
  - See Piazza for rooms
- Homework 6: Probability and Bayes Nets
  - Due Monday 10/17 at 11:59pm
- Project 3: MDPs and Reinforcement Learning
  - Due Friday 10/14 at 5pm

# Our Status in CS188

- We're done with Part I: Search and Planning!
- We've seen how AI methods can solve problems in:
  - Search
  - Constraint Satisfaction Problems
  - Games
  - Markov Decision Problems
  - Reinforcement Learning




# Our Status in CS188

Next up:

- Part II: Probabilistic Reasoning

- Diagnosis
- Speech recognition
- Tracking objects
- Robot mapping
- Genetics
- Error correcting codes
- ... lots more!

- Part III: Machine Learning



# CS 188: Artificial Intelligence

## Probability Review



Instructors: Adam Janin, Josh Hug -- University of California, Berkeley

# Recall (?): Inference with Bayes' Rule



- Example: Diagnostic probability from causal probability:

$$P(\text{cause}|\text{effect}) = \frac{P(\text{effect}|\text{cause})P(\text{cause})}{P(\text{effect})}$$

- Example:


- M: meningitis, S: stiff neck

$$\left. \begin{aligned} P(+m) &= 0.0001 \\ P(+s | +m) &= 0.8 \\ P(+s | -m) &= 0.01 \end{aligned} \right\} \begin{array}{l} \text{Example} \\ \text{gives} \end{array}$$

- From the information above, what is the chance that you have meningitis assuming your neck is stiff, i.e.  $P(+m | +s)$ 
  - In theory, you learned all this in CS70, and we'll be reviewing it today. Give it a shot and see how you do.


# Today

- Probability
  - Random Variables
  - Joint and Marginal Distributions
  - Conditional Distribution
  - Product Rule, Chain Rule, Bayes' Rule
  - Inference
  - Independence
- You'll need all this stuff A LOT for the next few weeks, so make sure you go over it now!



# Inference in Ghostbusters

- A ghost is in the grid somewhere
- Sensor readings tell how close a square is to the ghost
  - On the ghost: red
  - 1 or 2 away: orange
  - 3 or 4 away: yellow
  - 5+ away: green



- Sensors are noisy, but we know  $P(\text{Color} \mid \text{Distance})$

| $P(\text{red} \mid 3)$ | $P(\text{orange} \mid 3)$ | $P(\text{yellow} \mid 3)$ | $P(\text{green} \mid 3)$ |
|------------------------|---------------------------|---------------------------|--------------------------|
| 0.05                   | 0.15                      | 0.5                       | 0.3                      |

# Demo Ghostbusters – No probability

---



# Uncertainty

- General situation:
  - **Observed variables (evidence):** Agent knows certain things about the state of the world (e.g., sensor readings or symptoms)
  - **Unobserved variables:** Agent needs to reason about other aspects (e.g. where an object is or what disease is present)
  - **Model:** Agent knows something about how the known variables relate to the unknown variables
- Probabilistic reasoning gives us a framework for managing our beliefs and knowledge


|      |      |      |
|------|------|------|
| 0.11 | 0.11 | 0.11 |
| 0.11 | 0.11 | 0.11 |
| 0.11 | 0.11 | 0.11 |

|       |      |      |
|-------|------|------|
| 0.17  | 0.10 | 0.10 |
| 0.09  | 0.17 | 0.10 |
| <0.01 | 0.09 | 0.17 |

|       |       |      |
|-------|-------|------|
| <0.01 | <0.01 | 0.03 |
| <0.01 | 0.05  | 0.05 |
| <0.01 | 0.05  | 0.81 |

# Random Variables


- A random variable is some aspect of the world about which we (may) have uncertainty
  - $R$  = Is it raining?
  - $T$  = Is it hot or cold?
  - $D$  = How long will it take to drive to work?
  - $L$  = Where is the ghost?
- We denote random variables with capital letters
- Like variables in a CSP, random variables have domains
  - $R$  in  $\{\text{true}, \text{false}\}$  (often write as  $\{+r, -r\}$ )
  - $T$  in  $\{\text{hot}, \text{cold}\}$
  - $D$  in  $[0, \infty)$
  - $L$  in possible locations, maybe  $\{(0,0), (0,1), \dots\}$



# Probability Distributions

- Associate a probability with each value


- Temperature:



$P(T)$

| T    | P   |
|------|-----|
| hot  | 0.5 |
| cold | 0.5 |

- Weather:



$P(W)$

| W      | P   |
|--------|-----|
| sun    | 0.6 |
| rain   | 0.1 |
| fog    | 0.3 |
| meteor | 0.0 |

# Probability Distributions

- Unobserved random variables have distributions

$P(T)$

| T    | P   |
|------|-----|
| hot  | 0.5 |
| cold | 0.5 |

$P(W)$

| W      | P   |
|--------|-----|
| sun    | 0.6 |
| rain   | 0.1 |
| fog    | 0.3 |
| meteor | 0.0 |

Shorthand notation:

$$P(\text{hot}) = P(T = \text{hot}),$$

$$P(\text{cold}) = P(T = \text{cold}),$$

$$P(\text{rain}) = P(W = \text{rain}),$$

...

OK if all domain entries are unique

- A distribution is a TABLE of probabilities of values
- A probability (lower case value) is a single number

$$P(W = \text{rain}) = 0.1$$

- Must have:  $\forall x \ P(X = x) \geq 0$  and  $\sum_x P(X = x) = 1$

# Joint Distributions

- A *joint distribution* over a set of random variables:  $X_1, X_2, \dots, X_n$  specifies a real number for each assignment (or *outcome*):

$$P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n)$$

$$P(x_1, x_2, \dots, x_n)$$

- Must obey:  $P(x_1, x_2, \dots, x_n) \geq 0$

$$\sum_{(x_1, x_2, \dots, x_n)} P(x_1, x_2, \dots, x_n) = 1$$

$$P(T, W)$$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

- Size of distribution if  $n$  variables with domain sizes  $d$ ?
  - For all but the smallest distributions, impractical to write out!

# Probabilistic Models



- A probabilistic model is a joint distribution over a set of random variables
- Probabilistic models:
  - (Random) variables with domains
  - Assignments are called *outcomes*
  - Joint distributions: say whether assignments (outcomes) are likely
  - *Normalized*: sum to 1.0
  - Ideally: only certain variables directly interact
- Constraint satisfaction problems:
  - Variables with domains
  - Constraints: state whether assignments are possible
  - Ideally: only certain variables directly interact

Distribution over T,W

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

Constraint over T,W

| T    | W    | P |
|------|------|---|
| hot  | sun  | T |
| hot  | rain | F |
| cold | sun  | F |
| cold | rain | T |



# Events

- An *event* is a set  $E$  of outcomes

$$P(E) = \sum_{(x_1 \dots x_n) \in E} P(x_1 \dots x_n)$$

- From a joint distribution, we can calculate the probability of any event
  - Probability that it's hot AND sunny?  $P(\text{hot}, \text{sun})$
  - Probability that it's hot?  $P(\text{hot})$
  - Probability that it's hot OR sunny?  $P(\text{hot or sun})$
- Typically, the events we care about are *partial assignments*, like  $P(T=\text{hot})$

$P(T, W)$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

# Events



- $P(+\text{dog}, +\text{cat})$  ?

$P(D, C)$

- $P(+\text{dog})$  ?


- $P(-\text{cat OR } +\text{dog})$  ?

| D    | C    | P   |
|------|------|-----|
| +dog | +cat | 0.2 |
| +dog | -cat | 0.3 |
| -dog | +cat | 0.4 |
| -dog | -cat | 0.1 |



# Marginal Distributions

- Marginal distributions are sub-tables which eliminate variables
- Marginalization (summing out): Combine collapsed rows by adding



$P(T, W)$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |



$$P(t) = \sum_s P(t, s)$$

$P(T)$

| T    | P   |
|------|-----|
| hot  | 0.5 |
| cold | 0.5 |



$$P(s) = \sum_t P(t, s)$$

$P(W)$

| W    | P   |
|------|-----|
| sun  | 0.6 |
| rain | 0.4 |

$$P(X_1 = x_1) = \sum_{x_2} P(X_1 = x_1, X_2 = x_2)$$

# Marginal Distributions



$P(D, C)$

| D    | C    | P   |
|------|------|-----|
| +dog | +cat | 0.2 |
| +dog | -cat | 0.3 |
| -dog | +cat | 0.4 |
| -dog | -cat | 0.1 |

$$P(d) = \sum_c P(d, c)$$


$P(D)$

| D    | P |
|------|---|
| +dog |   |
| -dog |   |

$P(C)$

| C    | P |
|------|---|
| +cat |   |
| -cat |   |


$$P(c) = \sum_d P(d, c)$$



# Conditional Probabilities

- A simple relation between joint and conditional probabilities
  - In fact, this is taken as the *definition* of a conditional probability

$$P(a|b) = \frac{P(a, b)}{P(b)}$$



$P(T, W)$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

$$P(W = s|T = c) = \frac{P(W = s, T = c)}{P(T = c)} = \frac{0.2}{0.5} = 0.4$$

$$\begin{aligned} &= P(W = s, T = c) + P(W = r, T = c) \\ &= 0.2 + 0.3 = 0.5 \end{aligned}$$

# Conditional Probabilities



$P(D, C)$

| D    | C    | P   |
|------|------|-----|
| +dog | +cat | 0.2 |
| +dog | -cat | 0.3 |
| -dog | +cat | 0.4 |
| -dog | -cat | 0.1 |

■  $P(+dog \mid +cat) ?$

■  $P(-dog \mid +cat) ?$

■  $P(-cat \mid +dog) ?$

# Conditional Distributions

- Conditional distributions are probability distributions over some variables given fixed values of others

Conditional Distributions

|          |                        |     |
|----------|------------------------|-----|
| $P(W T)$ | $P(W T = \text{hot})$  |     |
|          | W                      | P   |
|          | sun                    | 0.8 |
|          | rain                   | 0.2 |
|          | $P(W T = \text{cold})$ |     |
|          | W                      | P   |
|          | sun                    | 0.4 |
|          | rain                   | 0.6 |

Joint Distribution

|           |      |     |
|-----------|------|-----|
| $P(T, W)$ |      |     |
| T         | W    | P   |
| hot       | sun  | 0.4 |
| hot       | rain | 0.1 |
| cold      | sun  | 0.2 |
| cold      | rain | 0.3 |

# Conditional Distributions (Tedious)

$P(T, W)$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

$$\begin{aligned}P(W = s|T = c) &= \frac{P(W = s, T = c)}{P(T = c)} \\&= \frac{P(W = s, T = c)}{P(W = s, T = c) + P(W = r, T = c)} \\&= \frac{0.2}{0.2 + 0.3} = 0.4\end{aligned}$$



$P(W|T = c)$

| W    | P   |
|------|-----|
| sun  | 0.4 |
| rain | 0.6 |

$$\begin{aligned}P(W = r|T = c) &= \frac{P(W = r, T = c)}{P(T = c)} \\&= \frac{P(W = r, T = c)}{P(W = s, T = c) + P(W = r, T = c)} \\&= \frac{0.3}{0.2 + 0.3} = 0.6\end{aligned}$$

# Conditional Distributions via Normalization Trick

$$\begin{aligned}P(W = s|T = c) &= \frac{P(W = s, T = c)}{P(T = c)} \\&= \frac{P(W = s, T = c)}{P(W = s, T = c) + P(W = r, T = c)} \\&= \frac{0.2}{0.2 + 0.3} = 0.4\end{aligned}$$

$P(T, W)$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

**SELECT** the joint probabilities matching the evidence



$P(c, W)$

| T    | W    | P   |
|------|------|-----|
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

**NORMALIZE** the selection (make it sum to one)



$P(W|T = c)$

| W    | P   |
|------|-----|
| sun  | 0.4 |
| rain | 0.6 |

$$\begin{aligned}P(W = r|T = c) &= \frac{P(W = r, T = c)}{P(T = c)} \\&= \frac{P(W = r, T = c)}{P(W = s, T = c) + P(W = r, T = c)} \\&= \frac{0.3}{0.2 + 0.3} = 0.6\end{aligned}$$

# Normalization Trick

$P(T, W)$

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

**SELECT** the joint probabilities matching the evidence



$P(c, W)$

| T    | W    | P   |
|------|------|-----|
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

**NORMALIZE** the selection (make it sum to one)



$P(W|T = c)$

| W    | P   |
|------|-----|
| sun  | 0.4 |
| rain | 0.6 |

- Why does this work? Sum of selection is  $P(\text{evidence})$ ! ( $P(T=c)$ , here)

$$P(x_1|x_2) = \frac{P(x_1, x_2)}{P(x_2)} = \frac{P(x_1, x_2)}{\sum_{x_1} P(x_1, x_2)}$$



# Normalization Trick



- $P(D \mid C=-\text{cat})$  ?

$P(D, C)$

| D    | Y    | P   |
|------|------|-----|
| +dog | +cat | 0.2 |
| +dog | -cat | 0.3 |
| -dog | +cat | 0.4 |
| -dog | -cat | 0.1 |

**SELECT** the joint  
probabilities  
matching the  
evidence



**NORMALIZE** the  
selection  
(make it sum to one)



# To Normalize

- (Dictionary) To bring or restore to a normal condition

All entries sum to ONE

- Procedure:

- Step 1: Compute  $Z$  = sum over all entries
- Step 2: Divide every entry by  $Z$

- Example 1

| W    | P   |
|------|-----|
| sun  | 0.2 |
| rain | 0.3 |

Normalize  
Z = 0.5

| W    | P   |
|------|-----|
| sun  | 0.4 |
| rain | 0.6 |

- Example 2

| T    | W    | P  |
|------|------|----|
| hot  | sun  | 20 |
| hot  | rain | 5  |
| cold | sun  | 10 |
| cold | rain | 15 |

Normalize  
Z = 50

| T    | W    | P   |
|------|------|-----|
| hot  | sun  | 0.4 |
| hot  | rain | 0.1 |
| cold | sun  | 0.2 |
| cold | rain | 0.3 |

# Probabilistic Inference

- Probabilistic inference: compute a desired probability from other known probabilities (e.g. conditional from joint)
- We generally compute conditional probabilities
  - $P(\text{on time} \mid \text{no reported accidents}) = 0.90$
  - These represent the agent's *beliefs* given the evidence
- Probabilities change with new evidence:
  - $P(\text{on time} \mid \text{no accidents, 5 a.m.}) = 0.95$
  - $P(\text{on time} \mid \text{no accidents, 5 a.m., raining}) = 0.80$
  - Observing new evidence causes *beliefs to be updated*



# Inference by Enumeration

- General case:


- Evidence variables:  $E_1 \dots E_k = e_1 \dots e_k$
  - Query\* variable:  $Q$
  - Hidden variables:  $H_1 \dots H_r$
- $$\left. \begin{array}{l} E_1 \dots E_k = e_1 \dots e_k \\ Q \\ H_1 \dots H_r \end{array} \right\} \begin{array}{l} X_1, X_2, \dots X_n \\ \text{All variables} \end{array}$$

- We want:

*\* Works fine with multiple query variables, too*

$$P(Q|e_1 \dots e_k)$$

- Step 1: Select the entries consistent with the evidence




| x  | P(x) |
|----|------|
| -3 | 0.05 |
| -1 | 0.25 |
| 0  | 0.07 |
| 1  | 0.2  |
| 5  | 0.01 |

2

0.15

- Step 2: Sum out H to get joint of Query and evidence



$$P(Q, e_1 \dots e_k) = \sum_{h_1 \dots h_r} P(Q, \underbrace{h_1 \dots h_r}_{X_1, X_2, \dots X_n}, e_1 \dots e_k)$$

- Step 3: Normalize

$$\times \frac{1}{Z}$$

$$Z = \sum_q P(Q, e_1 \dots e_k)$$

$$P(Q|e_1 \dots e_k) = \frac{1}{Z} P(Q, e_1 \dots e_k)$$

# Inference by Enumeration

- $P(W)$ ?
- $P(W \mid \text{winter})$ ?
- $P(W \mid \text{winter, hot})$ ?

| S      | T    | W    | P    |
|--------|------|------|------|
| summer | hot  | sun  | 0.30 |
| summer | hot  | rain | 0.05 |
| summer | cold | sun  | 0.10 |
| summer | cold | rain | 0.05 |
| winter | hot  | sun  | 0.10 |
| winter | hot  | rain | 0.05 |
| winter | cold | sun  | 0.15 |
| winter | cold | rain | 0.20 |

# Inference by Enumeration


---

- Obvious problems:
  - Worst-case time complexity  $O(d^n)$
  - Space complexity  $O(d^n)$  to store the joint distribution

# The Product Rule

- Sometimes have conditional distributions but want the joint

$$P(y)P(x|y) = P(x, y) \quad \longleftrightarrow \quad P(x|y) = \frac{P(x, y)}{P(y)}$$



# The Product Rule

$$P(y)P(x|y) = P(x, y)$$


- Example:

$P(W)$

| R    | P   |
|------|-----|
| sun  | 0.8 |
| rain | 0.2 |

$P(D|W)$

| D   | W    | P   |
|-----|------|-----|
| wet | sun  | 0.1 |
| dry | sun  | 0.9 |
| wet | rain | 0.7 |
| dry | rain | 0.3 |



$P(D, W)$

| D   | W    | P |
|-----|------|---|
| wet | sun  |   |
| dry | sun  |   |
| wet | rain |   |
| dry | rain |   |



# The Chain Rule

---


- More generally, can always write any joint distribution as an incremental product of conditional distributions

$$P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_1, x_2)$$

- By induction, we can chain this along infinitely many times.

$$P(x_1, x_2, \dots x_n) = \prod_i P(x_i|x_1 \dots x_{i-1})$$

# Bayes Rule



# Bayes' Rule

- Two ways to factor a joint distribution over two variables:

$$P(x, y) = P(x|y)P(y) = P(y|x)P(x)$$

That's my rule!

- Dividing, we get:

$$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$$

- Why is this at all helpful?

- Lets us build one conditional from its reverse
- Often one conditional is tricky but the other one is simple
- Foundation of many systems we'll see later (e.g. ASR, MT)

- In the running for most important AI equation!



# Bayes Rule Example

- Given:

$P(A)$

| A      | P    |
|--------|------|
| wolf   | 0.01 |
| sister | 0.99 |

$P(S|A)$

| S     | A      | P             |
|-------|--------|---------------|
| howl  | sister | 0.05          |
| hello | sister | 0.95          |
| howl  | wolf   | $1 - 10^{-6}$ |
| hello | Wolf   | $10^{-6}$     |

$$P(x|y) = \frac{P(y|x)}{P(y)}P(x)$$

- What is  $P(\text{sister} \mid \text{howl})$  ?

# Bayes Rule



- Given:

$P(W)$

| R    | P   |
|------|-----|
| sun  | 0.8 |
| rain | 0.2 |

$P(D|W)$

| D   | W    | P   |
|-----|------|-----|
| wet | sun  | 0.1 |
| dry | sun  | 0.9 |
| wet | rain | 0.7 |
| dry | rain | 0.3 |

$$P(x|y) = \frac{P(y|x)}{P(y)}P(x)$$

- What is  $P(W \mid \text{dry})$  ?

# Recall: Inference with Bayes' Rule



- Example: Diagnostic probability from causal probability:

$$P(\text{cause}|\text{effect}) = \frac{P(\text{effect}|\text{cause})P(\text{cause})}{P(\text{effect})}$$

- Example:

- M: meningitis, S: stiff neck

$$\left. \begin{array}{l} P(+m) = 0.0001 \\ P(+s | +m) = 0.8 \\ P(+s | -m) = 0.01 \end{array} \right\} \begin{array}{l} \text{Example} \\ \text{givens} \end{array}$$

- From the information above, what is the chance that you have meningitis assuming your neck is stiff, i.e.  $P(+m | +s)$ 
  - Will probably still be tricky!

# Inference with Bayes' Rule

- Example: Diagnostic probability from causal probability:

$$P(\text{cause}|\text{effect}) = \frac{P(\text{effect}|\text{cause})P(\text{cause})}{P(\text{effect})}$$

- Example:

- M: meningitis, S: stiff neck

$$\left. \begin{aligned} P(+m) &= 0.0001 \\ P(+s|+m) &= 0.8 \\ P(+s|-m) &= 0.01 \end{aligned} \right\} \text{Example gives}$$

$$P(+m|+s) = \frac{P(+s|+m)P(+m)}{P(+s)} = \frac{P(+s|+m)P(+m)}{P(+s|+m)P(+m) + P(+s|-m)P(-m)} = \frac{0.8 \times 0.0001}{0.8 \times 0.0001 + 0.01 \times 0.999}$$

- Note: posterior probability of meningitis still very small
- Note: you should still get stiff necks checked out! Why?

# Ghostbusters, Revisited

- Let's say we have two distributions:
  - **Prior distribution** over ghost location:  $P(G)$ 
    - Let's say this is uniform
  - Sensor reading model:  $P(R | G)$ 
    - Given: we know what our sensors do
    - $R$  = reading color measured at  $(1,1)$
    - E.g.  $P(R = \text{yellow} | G=(1,1)) = 0.1$
- We can calculate the **posterior distribution**  $P(G|r)$  over ghost locations given a reading using Bayes' rule:

$$P(g|r) \propto P(r|g)P(g)$$

|      |      |      |
|------|------|------|
| 0.11 | 0.11 | 0.11 |
| 0.11 | 0.11 | 0.11 |
| 0.11 | 0.11 | 0.11 |

|       |      |      |
|-------|------|------|
| 0.17  | 0.10 | 0.10 |
| 0.09  | 0.17 | 0.10 |
| <0.01 | 0.09 | 0.17 |



# Demo Ghostbusters with Probability

---

# Next Time: Bayes' Nets

---